

Analysis of Access to Emergency Funds in Sub-Saharan Countries– A Human Rights-Based Approach

Rose Porta^{1,†,*}, Alejandra Munoz Garcia^{1,†}, Margaret Bassney^{1,†}, Aushanae Haller^{1,†}

¹ Department of Statistical and Data Sciences, Smith College, Northampton MA;

* Correspondence: rporta@smith.edu

† These authors contributed equally to this work.

Abstract: Most people require access to emergency funds at least once in their life. These funds act as an important safety net in emergency cases. The purpose of our project is to predict access to emergency funds for adults in Sub-Saharan countries using a human rights-based approach to machine learning which centers equity, fairness, and impacts on humans over accuracy. Our analysis is based on the 2017 Global Findex Database which includes demographic as well as financial information for a sample of individuals within each country. We used a Decision Tree Classifier machine learning model implemented using Python to predict access to emergency funds with 68% accuracy. We assessed the fairness of our model with respect to gender using a variety of group and individual fairness metrics and evaluated the implications of each fairness metric with regard to our data and the goals of the analysis. We then implemented a variety of pre-processing, in-processing, and post-processing techniques in an attempt to minimize bias and maximize fairness. We have documented our analysis in a Jupyter notebook where this information can be made accessible to a broader undergraduate audience.

Keywords: Fairness; machine learning; access; emergency funds; financial; ethics; processing techniques; gender; economy; debiasing; decision tree; Sub-Saharan Africa; fairness metrics

1. Introduction

Science is often viewed as a way to offer trustworthy research backed solutions and answers. A lot of that research involves statistical methods performed on data. However, what happens when the data and statistical methods are not objective and trustworthy as is so often assumed? The conclusions drawn from the data are biased and unfair, most often towards minorities and protected classes of people. In order ensure that the outcomes of our analysis are equitable and positively impactful, our goal is to establish a human rights-based approach to machine learning analysis. A human right-based approach is characterized by awareness of context and consideration of the impacts the outcomes will have on humans through each stage of the data analysis process. More concretely, this process looks like (1) researching the context of the dataset before beginning any analysis; (2) performing exploratory data analysis to bring awareness to imbalances and relationships that exist in the data; (3) thoughtfully choosing an appropriate model based on the context of the data; (4) implementing fairness metrics to assess bias in the model; and (5) implementing de-biasing techniques to improve fairness.

We use a World Bank Global Findex data set which contains financial information about 35 Sub Saharan countries. Specifically, we create a model to predict access to emergency funds, defined as 1/20th of the GNI (gross national income) per capita for the country, then analyze the fairness of the

model. We focus on group and individual fairness metrics for the protected attribute sex. In addition we investigate the data set itself to understand where potential biases might have been implanted.

Data sets and algorithms have real world impacts on real people. The inherent bias in data sets can carry over into machine learning algorithms that are used to profile and categorize people [1,2]. Since data sets are not collected in a vacuum and often represent the discriminatory environments in which they are collected [3], we must find ways to make data sets and statistical methods more equitable. In this study we explore fairness methods that can be used to evaluate machine learning models. The “impossibility theorem” is the idea that not all fairness metrics can be satisfied at the same time [4]. Although fairness is complex and there are multiple approaches to make a model fair [5,6], it is important to continue to question how data and algorithms can be biased and how to mitigate that bias.

While there have been previous studies implementing fairness techniques in different contexts [7,8], we implement them in an exploratory context meant to teach how and when to use these techniques thus giving us more freedom to branch beyond a specific question while supporting previous work about the importance of these fairness metrics [3,9]. We analyse the data, data collection methods, prediction models, and the fairness metrics to assess how biased our data is and understand how we can de-bias when possible.

This approach will be used as the basis of a workshop meant to educate data science and computer science students on how to integrate ethics and fairness into their machine learning workflows. The analysis outlined in this paper is meant to serve as an example of an implementation of a human rights-based approach to machine learning. Thus, we focus more on explaining the process as it applies to this analysis more than the results of the analysis itself.

2. Data

Our data is derived from The World Bank in [The Global Findex Database](#), comprising the most comprehensive data sets on how adults save, borrow, make payments, and manage risk in different economies around the world. The data set was created to record various measures of financial equity and inclusion, with the intention that such information could reveal opportunities to expand access to financial services and to promote greater use of digital financial services for individuals who do not have a bank account. The data includes a sample of 1,000 randomly selected individuals from each of about 140 countries who completed interviews either over the phone or in-person depending on which made most sense in the given country. Each sample was drawn from the non-institutionalized population of adults aged 15 and over within the country. Intensive measures were taken to ensure that the samples were truly representative of the entire population in each country. For full details on the random sampling methodology, see the [data documentation](#).

There were several variables of interest in this dataset when creating models to predict access to emergency funds, including demographic and financial information. For this analysis, we are using only a subset of the data including countries in the Sub-Saharan region (35 countries total). Our data set includes 35,000 observations and 105 variables in total.

2.1. Demographics

2.1.1. gender

The variable *gender* distinguishes gender. There are 17,870 females and 17,130 males in the dataset.

2.1.2. education

The *education* variable corresponds to the highest level of education attained with ‘Primary’, ‘Secondary’ and ‘Tertiary’ being the three options. The distribution of education by gender plot (seen in Appendix Figure 8) shows us that there are more women with primary education, but more men with secondary or tertiary education. Overall, we can see that there are more men with higher education

than women. About 1,000 more men have received a secondary education and there is about double the amount of men with tertiary education compared to women, showing a clear disparity.

2.1.3. *economy*

The final demographic variable of interest is the *economy* variable that separates respondents by which country they live in. There are 35 different countries from Sub-Saharan Africa with exactly 1000 respondents from each. The countries included are Benin, Botswana, Burkina Faso, Cameroon, Central African Republic, Chad, Congo, Dem. Rep., Congo, Rep., Cote d'Ivoire, Ethiopia, Gabon, Ghana, Guinea, Kenya, Lesotho, Liberia, Madagascar, Malawi, Mali, Mauritania, Mauritius, Mozambique, Namibia, Niger, Nigeria, Rwanda, Senegal, Sierra Leone, South Africa, South Sudan, Tanzania, Togo, Uganda, Zambia, and Zimbabwe.

2.2. *Financial*

In the following, we will take a closer look at financial related variables that seem likely to have an impact to the access to emergency funds.

2.2.1. *account_fin*

The *account_fin* variable distinguishes those who have a financial account from those who do not. We found 11,970 respondents have a financial account while 23,030 respondents do not. This means about two thirds of individuals do not have an account. This is likely connected to the lack of access to emergency funds displayed above. If an individual does not have a financial account, we would expect they are less likely to have a source of emergency funds, as emergency funds are generally stored in an account.

2.2.2. *reason*

Those who do not have a financial account were asked why in the *reason* variable, that provides a list of possible reasons for not having a financial account (seen in Appendix Figure 9).

2.2.3. *emp_in*

Employment status was another financial variable of interest represented by *emp_in*, which asks whether or not the participant is in the workforce. It appears that 10,715 individuals were not in the workforce while 24,285 were, meaning about three-fourths of individuals are in the workforce.

2.2.4. *inc_q*

And lastly, we evaluated *inc_q*, which represents the relative socio-economic status of each respondent. Income quantiles divide a population into five income groups with 1 being the poorest and 5 being the richest such that about 20% of the population is in each group. The mean for all of the countries in the dataset is 3.241. This means that all the countries average out to be about middle class.

The majority of the data set has individuals within the richest quantile, Quantile 5.

2.3. *Emergency Funds*

To explore access to emergency funds in our dataset, which is our outcome variable of focus, we were interested three relevant variables:

2.3.1. *has_access*

The variable *has_access* directly asks participants if they have access to emergency funds, with "emergency funds" defined as 1/20th of the GNI (gross national income) per capita for the country. GNI per capita is the country's total income in a year/ the country's population size. For context, in the United States, "emergency funds" would be defined as about \$3,500 (U.S. GNI per capita as of




Figure 1. Distribution of Main Source of Emergency Funds

2021 = \$69,288 according to the [World Bank](#), and 1/20th of that is roughly \$3,500). In the Sub-Saharan region on average, the GNI per capita as of 2021 translated to U.S. dollars is \$1,578, and emergency funds (1/20th of that) would be defined as roughly \$80 ([source](#)).

The overall distribution of access to emergency funds showed 15,237 individuals had access to emergency funds while 18,867 did not. This indicates that over half of individuals represented in the data do not have access.

2.3.2. *main_source_funds*

We proceeded to explore the source of emergency funds using the *main_source_funds* variable, which provides a list of options for where participants receive their main source of emergency funds:

Figure 1 displays the overall distribution of the main source of emergency funds. Most of the individuals with access to emergency funds receive their funding from work, their family and friends, or their savings.

2.3.3. *Receive Wage Payments*

The “Receive Wage Payments” variable is related in that receiving wages is necessary in order to have money from work. Only 8,348 individuals receive wage payments while 26,241 do not. This suggests that receiving wage payments may be a key factor in determining access to emergency funds, yet relatively few individuals receive wages.




Figure 2. Percent of Participants by Country & Gender

2.3.4. *gender, economy and education* in relation to Emergency Funds

Finally, we sought to identify disparities in access to emergency funds by *gender, economy, and education*.

In the side-by-side barplots shown in figures 2 and 3, we can see that although only about 50% of the countries have a higher percentage of men represented in the questionnaire (left bar plot), in 75% of the countries more men have access to emergency funds than women (right bar plot).

Additionally, in figure 4, we can see the distribution of funds based on an individual's highest education level. 63% of people with only a primary education do not have access to emergency funds compared to 37% of people who do. These numbers are more evenly distributed for those with secondary education, with about 49% of people not having access to emergency funds, while 51% of people do have access. Finally, for those with a tertiary level of education we can see that about 72% of people have access to emergency funds while only 28% of that group does not have access. Overall, we can make the assumption that people with a higher level of education are more likely to have access to emergency funds.

Overall, our exploratory data analysis reveals that there are disparities in access to emergency funds based on gender and education level. Thus, it will be essential to take these imbalances into account when fitting our model. In this analysis, we choose to focus on maximizing equity in model predictions with respect to gender. However, if these predictions were being used to make real-world decisions that would impact people's lives, it would be necessary to ensure that bias based on education level is also minimized.




Figure 3. Percent of Participants with Access to Emergency Funds by Country & Gender




Figure 4. Access to Emergency Funds by Education Level

3. Methods

In this section, we describe the processes used for fitting and modifying our model. We detail our processes of data cleaning, selecting an appropriate model, choosing and implementing metrics to assess the model, and choosing and implementing de-biasing techniques to improve fairness.

3.1. Software

We conducted our analysis in a Google Colab Notebook primarily employing the Python libraries `pandas` and `numpy` for data cleaning and exploratory analysis, as well as `sklearn` and `aif360` for implementing machine learning methods, fairness metrics, and de-biasing techniques. The notebook detailing our [full analysis](#) is available via our public [GitHub repository](#).

3.2. Data Cleaning

Before fitting the model, we performed several pre-processing steps on the data in order to remove unnecessary or redundant information, address missing values, and ensure that the variables were coded such that they would be processed appropriately by the model. Some of these steps were performed after fitting the initial model, and some choices were made based on the impacts of those choices on the performance of the model.

First, we removed the arbitrary variables *economycode* (country code) and *regionwb* (region); *economycode* is essentially a duplicate of *economy* (country name), and *regionwb* is the same for all rows (value is Sub-Saharan) since this is the variable that we initially filtered by.

Next, we checked for variables with a high percentage of missing values. Several variables have many missing values because they are follow-up questions to a previous question that are only asked if the respondent gives a specific response for a previous question. We chose to drop all variables with more than 30% missing values (58 variables total) because we observed by running the model multiple times that variables with NA percentages above this threshold had no impact on the model accuracy or fairness. Thus, it made sense to remove them if their presence is negligible when included in the model.

Next, we checked for variables with high levels of redundancy (i.e. little variation), defining a high level of redundancy as 95% or more of the values being the same. The only variable with a high level of redundancy was *pay_online* (a binary variable indicating whether or not an individual has paid online for something), and we chose to remove it because removing it had no impact on the model accuracy and fairness.

The most common answer format of the survey was multiple choice with answer options “yes”, “no”, “don’t know”, and “refuse”. Based on our exploratory data analysis, it seems that for most questions, the numbers of “don’t know” and “refuse” responses are very low. Furthermore, these responses would not give us much useful information when implementing a predictive model. Thus, we chose to replace all “don’t know” and “refuse” values with NA values. We then replaced all remaining missing values (including missing values not removed previously and “don’t know” and “refuse” values) with the column mean rounded to the nearest whole number (i.e. the most frequent value if the variable is categorical).

Next, we re-coded the country variable into a variable with five categories based on the percentage of sampled individuals in the country who have access to emergency funds (1 = < 20% have access, 2 = between 20% and 40% have access, etc.). All other variables in our data set were coded such that the model could appropriately interpret them, so we did not have to do any additional re-coding. The majority of the variables are binary with 0 = no, 1 = yes, and the rest are either categorical variables with hierarchical categories such as income quantile and education level or continuous numerical variables such as age.

Finally, we created a cartesian product to combine two variables – one indicating whether or not an individual has a financial account and the other indicating whether or not the individual has saved

money in the past 12 months– in order to increase the accuracy of the model after fitting the initial model.

After this pre-processing, we have 42 predictor variables and 35,000 observations. A full list of the predictor variables included in our model can be found in the appendix.

3.3. Model Selection

Given that we are aiming to predict a binary outcome (possible or not possible to come up with emergency funds), our model needs to be a classification model. We tried two of the most common model types used for classification– logistic regression and decision tree classifier. We ultimately chose the decision tree classifier model over the logistic regression model because the baseline accuracy was higher (61% versus 55%). Furthermore, the decision tree model makes more intuitive sense given our data since most of our predictors are binary variables, and binary predictors fit well into the tree structure. We were able to further improve the accuracy of the decision tree model to 68% by using cross-validation to specify the max depth as 6. The decision tree model uses machine learning to predict outcomes by organizing the variables into a tree that branches off at each decision point based on the value of the variable at that point. The most influential variables are at the top, and the outcome variable is at the end of each branch. We split our data into 70% training, 30% testing because this is the standard train-test split used for machine learning algorithms [10].

3.4. Metrics

To assess fairness and accuracy in our model we explore 12 different metrics, 10 of which are fairness metrics. Fairness metrics can be split into group and individual metrics.

3.4.1. Basic Metrics

Accuracy.

Accuracy is a measure of how many classifications our model predicts correctly compared to all the predictions; i.e., the ratio of correctly predicted classifications to all the predictions. Accuracy cannot tell us if the predictions are equally correct across positives and negatives [11,12]. 55% of the people in our data set do not have access to emergency funds. As long as our model predicts negatives more than half the time, we can get a good accuracy. However, our model will lose the ability to accurately predict positives. It is important to consider accuracy along with precision and recall (both explained in following sections) so we can more fully understand how our model is classifying people.

$$Accuracy = (TruePositives + TrueNegatives) / (TruePositives + TrueNegatives + FalsePositives + FalseNegatives)$$

Precision.

Precision is a measure of how accurately a model predicts positive outcomes; i.e., the ratio of correctly predicted positives to all predicted positives. With high precision rates, we have low false positive rates [11,12].

$$Precision = truepositives / (truepositives + falsepositives)$$

Recall.

Recall is a measure of how accurately a model predicts negative outcomes; i.e., the ratio of correctly predicted negatives to all predicted negatives[12]

$$Recall = truenegatives / allnegatives)$$

3.4.2. Fairness Metrics

Fairness metrics are a way to assess machine learning algorithms for unwanted bias. Algorithms can classify people unfairly using data collected in a biased environment. When classifying people, it is important to take into account how these classifications can contribute to and reinforce discriminatory social systems. Thus, it is sometimes necessary to sacrifice accuracy in favor of fairness when using machine learning algorithms to make decisions impacting people [13, Kamiran and Calders [14], Menon and Williamson [15]]. It is often not possible to maximize fairness and accuracy at the same time because due to discriminatory systems built into society, protected attributes such as gender are often associated with the outcome variable. With regard to our dataset, the proportion females with access to emergency funds is substantially lower than that for men in the sample. Because of this, the model will be less likely to predict positive values for females. If the model were to hypothetically be used to determine which individuals to offer loans to, it would reinforce existing patriarchal systems to offer fewer loans to females. Thus, it would be beneficial to prioritize equitable access to loans by gender over accuracy. In other words, reality is often not fair, so if we only prioritize accuracy, we will continue to replicate the discriminatory systems that exist.

One approach to fairness in machine learning is “fairness by unawareness” meaning a model is blind to the sensitive attributes. Although it may seem an intuitive approach to simply remove the protected attribute from the data in order to make the algorithm unbiased, this is often not an effective approach to reduce bias. There are often variables that remain in the data that act as pseudo substitutes for the protected attribute [16]. For example, if race was excluded from the model but the variable zip code remained. Zip code can act as a stand in for race in regions where people are segregated by race. Another concrete example of this is the Amazon recruiter tool removed gender from their algorithm but it was still biased against women [17].

In our case we are focusing on the protected attribute *gender*. If we were to remove the gender variable from our data, it is possible that our model could still have gender bias given that there are other variables in the dataset that have a strong association with gender including education and several of the financial variables (as shown in our exploratory data analysis).

In this section, we describe the seven fairness metrics that we implemented. We chose these specific metrics based on the context of our data and the objectives of our analysis. The metrics that are most informative change depending on the context. For example, since we are focusing on between-group gender differences, we chose to implement more group fairness metrics than individual fairness metrics. More specifically, equal opportunity difference is the most informative fairness metric in the context of our analysis because it focuses on equity with respect to accurately predicted positive values. We want the model to predict positive values with equal accuracy for males and females, and we want to minimize false positives to ensure that folks who do not have access and thus need financial assistance get the help that they need (assuming a context where model predictions are used to determine to whom to provide financial assistance based on who does not have access). All of the fairness metrics that we implemented along with context around when they are most useful are explained in the following sub-sections.

Individual Fairness Metrics

Individual fairness metrics measure how similarly the model predicts for similar observations. Will two very similar people receive the same classification? Individual fairness metrics contradict group fairness metrics. When accounting for imbalanced predictions between groups, the within group fairness can suffer [5]. In the process of satisfying group metrics, two similar subjects only differing by sex, may be classified differently [16,18–20]. The only individual metrics we explore is general entropy error (defined in the following sub-section), as we are prioritizing group fairness.

General Entropy Error. This metric is an individual metric, and it computes fairness by computing the level of unfair benefit being assigned by the model. The metric defines “benefit”

as follows: for any individual in the testing data set, that individual has received a benefit if the model predicted the favorable outcome when the truth was that the individual did not have the favorable outcome (i.e. a false positive). Each individual in the data receives either a 2 (benefit, false positive), a 1 (no benefit, correct prediction), or 0 (no benefit, false negative). The metric then compares the benefit of each individual to the average accuracy and false positive level of the model. The “ideal” value is 0, and a higher number indicates a higher level of inequity in benefit among individuals. In other words, if many individuals have a benefit score that is far off from the average, that indicates that the model is unfairly benefiting some individuals and not others. This metric does not consider privileged versus unprivileged groups, and thus is not able to indicate whether or not the inequality in benefit is systematic in any particular way (i.e., it cannot tell whether males receive more benefit than females; it can only tell that some individuals receive higher benefits than others) [5,20].

This metric is important to our data because a false positive outcome (higher benefit) would mean that someone is predicted to having access to emergency funds when they do not. If there is a high general entropy error, there are many individuals whose need for emergency funds are being overlooked.

Group Fairness Metrics

Group fairness metrics ensure parity between privileged and unprivileged groups of a protected class. For example, for the protected class sex, the privileged group is men and the unprivileged group is women. Group fairness metrics measures how discriminatory the model classifies the unprivileged group [18–20]. The group metrics we explore include statistical parity difference, equal opportunity difference, disparate impact, precision score difference, general entropy difference, and conditional demographic parity (explained in the following sub-sections). Not all group fairness metrics can be satisfied at the same time. For example equal opportunity difference and statistical parity difference cannot be simultaneously accounted for [5].

Statistical Parity Difference. This metric computes the difference in percentages between the “privileged” and “non-privileged” group of individuals who were predicted to have the desired outcome. In this case, it is essentially

$(\% \text{ of females who were predicted to have access to emergency funds}) - (\% \text{ of males who were predicted to have access to emergency funds})$

The “ideal” value is 0 because if we define fairness as statistical parity, the goal would be for the percentages to be equal for both groups. If the value is negative, that means that the percentage of individuals with the positive outcome is higher for the privileged group (males), implying that the model is biased in favor of the privileged group. Conversely, if the value is positive, the model is biased in favor of the unprivileged group. The acceptable range in which the model is considered fair is between -0.1 to 0.1 (with percentages expressed as decimals, e.g. 0.1 = 10%). It is important to note that this metric is solely focused on making the percentage of *predicted* favorable outcomes equal across groups and does not take into account the accuracy of the predictions at all [5,20].

Relating to our data, this metric will tell us if our model predicts that men have more access to emergency funds than women. In our data, men in fact do have more access to emergency funds than women. 51% of men have access to emergency funds while only 38.2% of women have access to emergency funds. When using this metric to asses our model the interpretation depends on the context. If this model is being used to decide how to allocate emergency funds, we might not want to prioritize satisfying this metric. We are using this model in an educational and exploratory manner, so we will use techniques to account for this metric.

Equal Opportunity Difference. This metric is similar to statistical parity in that it is also a group fairness metric, but it is different in that it takes into account accuracy of the model in addition to equalizing outcomes across groups. Instead of measuring the simple differences in percentages

between groups of individuals with the (predicted) positive outcome, it measures the difference in percentages of *accurately identified* individuals with positive outcomes (i.e. true positives). Essentially, the calculation is the same as for statistical parity, but only taking into account true positives for each group. Again, the “ideal” value is 0 with negative values indicating bias in favor of the privileged group, and the fairness range is -0.1 to 0.1 [20].

This metric helps us answer if our model predicts positives with more accuracy for men than women; i.e., are men more accurately predicted to have access to emergency funds than women?

Disparate Impact. The disparate impact metric measures the proportion of positive outcomes between an unprivileged group and a privileged group. It is usually assessed when predicting an outcome that disproportionately affects a sub population. For example, hiring more men than women as construction workers on the basis of height and strength. For this case we want to know the proportion of females that are categorized as having access to emergency funds versus males who are categorized as having access to emergency funds. The standard for satisfying this metric is that the unprivileged group must receive a positive outcome at a ratio of 4:5 to the privileged group. As long as females are classified as having access to emergency funds no less than around 80% of the time males are categorized as having access to emergency funds, then our model satisfies this metric [20]. This metric is similar to statistical parity except it measures a ratio which can be useful for legal purposes. $P(\hat{Y} = \text{unprivilegedPositivePredicted}) / P(\hat{Y} = \text{privilegedPositivePredicted})$

A similar problem arises when assessing this metric as statistical parity. In reality women have less access to emergency funds than men. If we manipulate our model to satisfy this metric, we will falsely predict that women have access to emergency funds when they do not. This could be more harmful than not satisfying this metric.

Conditional Demographic Disparity. Statistical parity difference and equal opportunity difference both measure positive outcomes. The conditional demographic disparity measures negative outcomes. Demographic Disparity is a metric that examines how disadvantaged groups compare to advantaged groups for negative outcomes from the model. This metric checks if a subpopulation is classified with a negative outcome more than a positive outcome; are females classified as not having access to emergency funds more often than men? Looking at the entire data set, women have less access than men to emergency funds in reality, and thus predicting more negative outcomes for women than men is not necessarily a bad thing. We want to know if someone does not have access to emergency funds so that they can potentially be helped.

Sometimes when we split data into categories we can find patterns that do not exist when the data is combined. This is called Simpson’s Paradox [19]. We can see this in our confusion matrices. When our data is split by gender there are different prediction rates than the entire model. The true negative rates are heavily weighted by females, and the true positive rates are weighted towards males. Is this the Simpson’s paradox, or does gender split the data into different distributions? The Conditional Demographic Disparity metric accounts for the Simpson’s paradox to confirm true differences or no differences in negative outcomes in the model.

The range of scores for this metric is from -1 to 1. In general a positive value means that the model is more unfair towards the unprivileged group. A value of zero is ideal. In our case, if our model were to predict an equal proportion of negative and positive outcomes for men and women, our model would realistically be unfair to women. Women do have less access to emergency funds, and predicting that men and women equally do not have access to emergency funds might put women at a greater disadvantage if a relief program were to be put in place. However, if assessing financial stability between men and women we would be more concerned with satisfying this metric.

General Entropy Error Difference. The general entropy error cannot tell whether males receive more benefit than females so we calculate the general entropy error difference between males and

females. Do men “benefit” from our model more; i.e., does our model predict more accurate and false positives for men than women? Again the interpretation of this metric will depend on the context. A higher score is given for false positives, but this means a group of people are being predicted to have emergency funds when they do not. It is not necessarily a good thing for any group to “benefit” from our model. A value of 0 represents no difference.

Between Group Generalized Entropy Error. We explored generalized entropy error and how it differs for males and females. Using the between group generalized error metric we will be able to see if the between group unfairness or the individual unfairness dominates. Is there truly a difference in generalized entropy error between men and women, or is the general entropy error not due to gender inequality? Is our model unfairly benefiting individuals based on sub populations or is the inequity equal between groups and differs at the individual level?

We do not want generalized entropy error in our model, but it would be better to have it at the individual level than the group level. We do not want either men or women to have more generalized entropy error than the other. If the error is equally within the groups, then both men and women are at a similar “benefit” to each other.

3.5. De-biasing Techniques

To account for any unfairness we find in the model we can use pre, in, and post processing techniques. These techniques restructure the data and reclassify observations in order to satisfy these fairness metrics. We implement several de-biasing techniques for educational purposes in order to demonstrate how they work and what their impacts are. However, it is important to consider that de-biasing techniques are not always necessary and are always context-dependent. If we are using model predictions to make choices that impact people’s lives, it generally makes sense to implement de-biasing techniques to ensure that all protected groups are being treated equally. For example, if we are using model predictions to decide to whom to offer loans, it makes sense to ensure that females and males receive loans at equitable rates. However, if we are seeking to use our model predictions to gain insight the reality of society (which is unfair), it may make more sense to prioritize accuracy and not implement de-biasing techniques. For example, if we are seeking to know who actually has access to emergency funds so that we can help those who do not, it would actually be beneficial for our model to predict more negative values for women; women do have less access to emergency funds in reality, so this would mean that more women who are in need would be receiving help. In this case, if we were to implement fairness metrics in order to make the rate of positive predictions equal for males and females, the model would be predicting more false positives, which would mean that many women would go without getting the help that they need.

3.5.1. Reweighting.

Reweighting is a pre-processing technique which assigns weights such that the protected attribute (gender) becomes statistically independent from the outcome variable (access to emergency funds). This means that after reweighting, knowing the gender of an individual does not provide any information about whether or not the individual has access. In mathematical terms, $P(\text{gender} = \text{male and access} = \text{yes}) = P(\text{gender} = \text{male}) * P(\text{access} = \text{yes})$, and this equality holds true for all gender-access combinations.

3.5.2. Exponentiated Gradient Reduction.

The exponentiation gradient reduction is an in-processing optimization approach. This processor aims to optimize both accuracy and fairness focusing on demographic parity and equalized odds. The algorithm this processor uses considers randomized classifiers and cost constraints to find the optimal classifier that satisfies fairness restraints without losing too much accuracy [21].

3.5.3. Grid search Reduction.

Grid search reduction uses the cost constraint lamda to find a balance between fairness and accuracy. This processor searches over a grid of lamda values until the best value is found. This value is used in the classifier to satisfy fairness and maximize accuracy. The grid search reduction is useful for binary sensitive attributes and fairness metrics with minimal constraints like demographic parity and equalized odds [21,22]

3.5.4. Calibrated Equalized Odds.

Calibrated equalized odds uses a post-processing technique that re-classifies values to satisfy the equalized odds metric while keeping the classifier calibrated. A classifier is calibrated if the proportions of positive and negative outcomes in the data match the probabilities produced by the model. We want the calibration to hold across groups such as male and female. This processor aims to satisfy an equalized cost constraint while maintaining calibration [23].

3.5.5. Reject Option Classifier.

The reject option classifier is a post-processor that aims to reduce discriminatory classifications based on the sensitive attribute. In our case we aim to find a balance for predictions between males and females. This classifier will relabel observations in a way that reduces discrimination. More males will be relabeled with the unfavorable outcome and more females will be relabeled with the favorable outcome [24].

3.5.6. Meta Fair Classifier.

The meta fair classifier creates a new estimator but includes a reweighing pre-processing step [25]. This classifier should be used as part of a pipeline of steps. We must create a binary label data set. This means that the data includes either a 1 representing access to emergency funds or 0 for no access to emergency funds. This classifier aims to transform the data in a way that will satisfy as many fairness metrics as possible [21].

3.6. Results

3.6.1. Decision Tree Model (Figure 5)

We fit a decision tree classifier model with the outcome variable *has_access* using 42 predictor variables (listed in appendix) and 24,000 data points in the training set (70% of 35,000). The decision tree model organizes the variables into a tree branching off at each decision point based on the value of the variables at that point. The most influential variables are at the top, and the outcome variable is at the end of each branch. In our decision tree, the most influential variable was *sent_remittances*; *sent_remittances* is a binary variable that asks participants, "In the PAST 12 MONTHS, have you, personally, GIVEN or SENT money to a relative or friend living in a different city or area INSIDE (country where survey takes place)? This can be money you brought yourself or sent in some other way." It makes sense that this is the most influential variable in relation to access to emergency funds, because if you have the privilege to send spare money to loved ones, you most likely have extra money saved for yourself as well.

Depending on the response to *sent_remittances* the tree branches into two, considering participants' responses to the next two influential variables: *access_cat* and *account_saved*; *access_cat* is a variable that categorizes the 35 sub-saharan countries by the percent of access to emergency funds according to its participants ranging between 1-5, with 1 being 0%-20% of participants in that country having access to emergency funds, and 5 being 80%-100% of participants having access to emergency funds. Depending on the country the participant lives in, they will be assigned to a number between 1-5 for *access_cat*. On the other hand, *account_saved* is a variable that categorizes participants based on their personal

Decision Tree Classifier to Predict Access to Emergency Funds

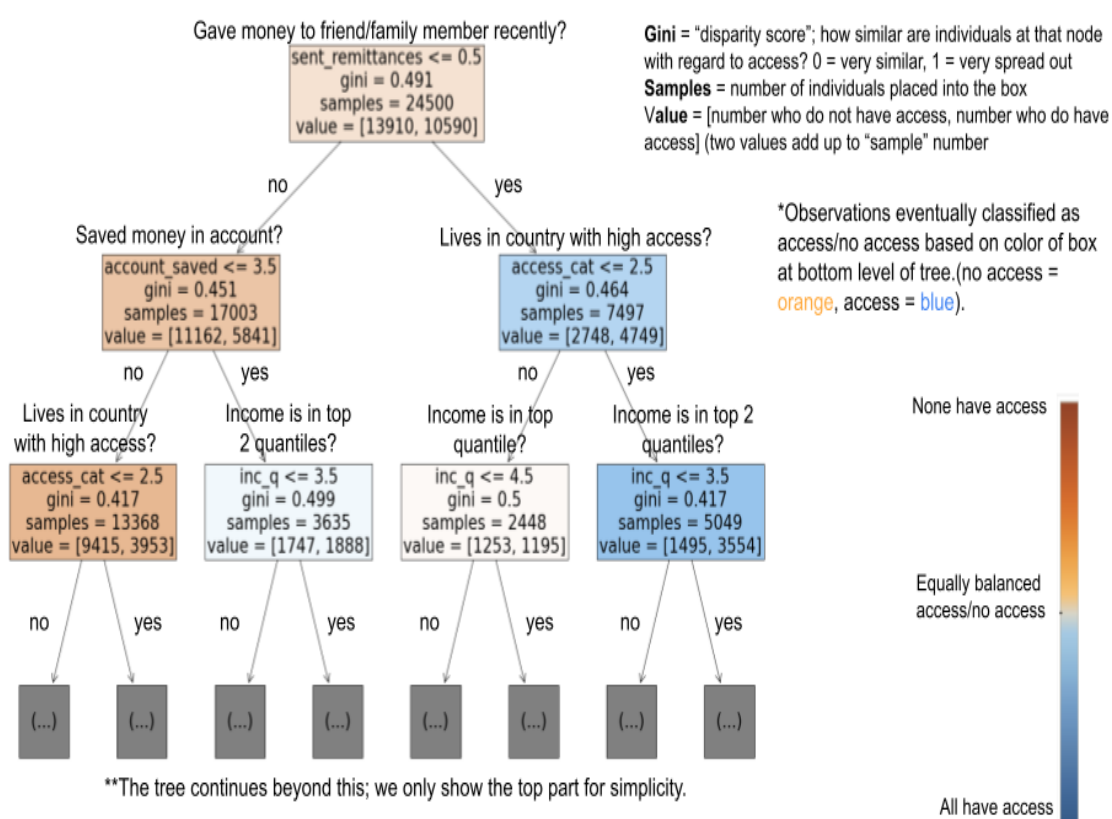


Figure 5. Decision Tree Classifier Model




Figure 6. Confusion Matrices by Gender

finances ranging from 1-4, with 1 representing having no financial account and no money saved, and 4 representing having a financial account and money saved. Being assigned 2 or 3 in *account_saved*, means that the respondent either only has a financial account or only has money saved, but not both.

The decision tree then splits according to participants' responses all the way down the tree, until finally it can predict whether or not based on the responses the participant has access to emergency funds.

A visualization of the decision tree is shown in figure 5. It is notable that as the tree branches off to the left, the boxes become darker and darker orange, meaning that a higher percentage of individuals in those boxes do not have access. As the tree branches to the right, the boxes become increasingly darker blue, indicating that a higher percentage of individuals have access. Thus, the further toward the right that an observation moves, the more likely that observation is to be classified into a positive outcome (has access). At the bottom of the tree, each observation is classified based on the color of the final box that it is placed into (blue = has access, orange = does not have access). The "samples" label in each box indicates how many individuals have been placed into that box. The "values" label indicates how many of those individuals do not have access and do have access respectively.

3.6.2. Confusion Matrices by Gender (Figure 6)

When we create two separate confusion matrices separated by gender, we can see that the model is fairly good at predicting true negatives for both groups, but is significantly better at predicting true positives for males as compared to females. The false positive and false negative error rates are also slightly higher for females as compared to males, yet the rates are pretty similar for both models.

This may be an example of the model reinforcing biases that are present in the data because we saw from our exploratory data analysis that many fewer females than males have access to emergency funds, so if there are not many females that have true positive values recorded in the data, the model is likely to be less accurate in predicting positive outcomes for females. This could have discriminatory impacts in the case that this model is (hypothetically) used to determine whether or not to provide an individual with a loan based on whether or not they have access to emergency funds (i.e., if they do have access, they are likely to pay back the loan, so they will get the loan, and if they do not have access, they will not get the loan). If very few females in the training data have access, and the model is thus less likely to predict a positive outcome for females as we see here, this will result in fewer females receiving loans. This will result in a self-reinforcing cycle of continued discrimination because if the model continues to predict that females should not get loans, fewer females will get loans, which means fewer females would have the opportunity to pay back loans, and so on.

On the flip side, if the results of the model are (hypothetically) being used to offer support to individuals who are in financial hardship and do not have access to emergency funds, the accuracy in predicting negative values would be more important, and the fact that the model is less accurate in predicting true positives would be less of a concern. This is because in this case, it would be a better outcome for an individual to receive support when they do not need it (false negative) as compared to the outcome that an individual does not receive support when they do need it (false positive). However, we would still want the model to be as accurate as possible on both ends in order to ensure that we are allocating resources most directly to those who need it the most.

This comparison of confusion matrices also demonstrates the importance of considering the balance of true negatives to true positives in addition to just the error rates. When looking at just the error rates, they seem almost equal between the two groups, yet when we consider the balance of true negatives to true positives, we can see that the model is much worse at predicting true positives for females. This furthermore highlights the importance of trying out various methods and metrics for assessing fairness because two different metrics can tell two completely different stories about the degree of fairness present in the model.

3.6.3. Fairness Metrics Summary Table (Figure 7)

The fairness metrics table shown in figure 7 summarizes the changes in fairness metrics in response to the (Pre, In, & Post) Processing Methods that we have applied to our model. Each column represents a metric, and each row represents a processing techniques are in the rows. The “objective” row (first row) shows the “ideal” values for each metric, while the “baseline” row (second row) shows the values for our baseline model without any processing techniques applied.

Baseline

Our baseline model demonstrated the second best model accuracy and precision at .684 and .682 respectively. It also has the closest general entropy error difference ratio to 0, at .0007, meaning that there is almost no difference in our baseline model between the benefits men and women have when it comes to having access to emergency funds.

Reweighting

Our reweighting technique demonstrated the highest model accuracy, at .689, meaning that our model predicts our classifications the most correctly when the reweighting technique is applied. Additionally, the reweighting technique also provides us with the best ratio for precision, at .683, indicating that applying this technique provides us with the best model for predicting positive outcomes accurately.

Fairness Metrics Table

| | technique | accuracy | precision | prec_diff | stat_par_diff | eq_opp_diff | gen_entr_error | disp_impact | gen_entr_diff | bet_grp_gen_entr | cond_dem_disparity |
|---|--------------------------------|----------|-----------|-----------|---------------|-------------|----------------|-------------|---------------|------------------|--------------------|
| 0 | objective | 1.000 | 1.000 | 0.000 | 0.000 | 0.000 | 0.000 | 1.000 | 0.000 | 0.000 | 0.000 |
| 1 | baseline | 0.684 | 0.682 | -0.103 | -0.161 | -0.170 | 0.269 | 0.625 | 0.007 | 0.000 | 0.003 |
| 2 | reweighting | 0.689 | 0.683 | -0.109 | -0.114 | -0.101 | 0.259 | 0.727 | -0.012 | 0.000 | 0.002 |
| 3 | meta fair classifier | 0.553 | 0.370 | -0.093 | 0.004 | 0.004 | 0.581 | 1.285 | -0.210 | 0.007 | -0.001 |
| 4 | exponential gradient reduction | 0.668 | 0.653 | -0.118 | -0.072 | -0.045 | 0.273 | 0.820 | -0.033 | 0.001 | 0.001 |
| 5 | grid search reduction | 0.668 | 0.673 | -0.143 | -0.053 | -0.041 | 0.297 | 0.847 | -0.038 | 0.001 | 0.001 |
| 6 | calibrated equalized odds | 0.681 | 0.662 | -0.113 | -0.193 | -0.217 | 0.254 | 0.596 | 0.027 | 0.001 | 0.003 |
| 7 | reject option classifier | 0.671 | 0.666 | -0.216 | 0.056 | 0.072 | 0.282 | 1.180 | -0.059 | 0.005 | -0.001 |

Red: Not within accepted range

Green: Within accepted range

Blue: individual metric

Figure 7. Fairness Metrics Summary

Meta Fair Classifier

The meta fair classifier provides us with an acceptable statistical parity difference and equal opportunity difference ratio of .004 for both, meaning that there is nearly no difference between the percentage of females who have access to emergency funds compared to the percentage of males who have access to emergency funds. However, because the ratio is positive, it indicates that the model is slightly biased in favor of women.

Both the statistical parity difference and equal opportunity difference values for the meta fair classifier were at .004, which is within the acceptable range. The statistical parity difference value indicates there is nearly no difference between the percentage of females who have access to emergency funds compared to the percentage of males who have access to emergency funds. The equal opportunity difference value indicates that there is nearly no difference between the accuracy in the model for women compared to the accuracy within the model for males. However, because these values are positive, they indicate that the model is slightly biased in favor of women.

Exponential Gradient Reduction & Grid Search Reduction

The exponential gradient reduction and grid search reduction techniques output approximately the same values for our fairness metrics. For the disparate impact metric, we had a ratio of about .847, which satisfies the acceptable range of $> .80$. This means that women are classified to have access to emergency funds about 84.7% of the time that men are categorized to have access to emergency funds. For our between group generalized entropy error metric, we have a value of .001, indicating that we have near equal levels of inequity for females and males. This means that there is not a large between-group disparity with regard to how much females versus males benefit from the model.

Calibrated Equalized Odds

Our calibrated equalized odds technique provided us a general entropy error difference of .254. A value of zero indicates that no group benefits from our model. This value (.254) indicates that women

benefit from our model and therefore are predicted to have more false positives, which means that women are projected to have access to emergency funds when they in reality do not.

Reject Option Classifier

The reject option classifier provides us with a value of -.001 for the conditional demographic disparity metric, meaning that there is close to an equal number of negatives and positives between groups. However, in our particular instance, this indicates an unfairness towards women, since we are aware that women have less access to emergency funds compared to men. This value indicates equity between our groups, which we know is false.

4. Discussion

In this study we examine a decision tree classifier for access to emergency funds in Sub-Saharan Africa, assess the fairness of this classifier, use processing techniques to increase the fairness of our model, and assess the fairness again after the processing techniques have been applied. Our initial decision tree model is unfair towards women in that it cannot predict true positives for women as well as it does for men. We apply six different processing techniques and analyze the results. We find that the processing techniques maintain most of the accuracy and improve fairness. In our case we specifically care that the model does not discriminate based on gender. This means we care most about the group fairness metrics and the processors that decrease gender discrimination in our model.

Most of the processors increased fairness for at least some of the fairness metrics. The calibrated equalized odds however, decreased fairness for all the metrics except the generalized entropy error. This would be a good processor to use if individual fairness was a major concern in the model. For group fairness, in our model, this processor is not ideal.

All processors besides calibrated equalized odds decreased the statistical parity difference resulting in our model predicting more women as having access to emergency funds. We must be careful when interpreting this result. At first glance it may seem that these techniques are increasing fairness, but we must consider how the other metrics change and what this means in a real-life context. Our processed model will classify more women as having access to emergency funds, but this could be dangerous if these women are realistically in need of financial support. To make sure our model is realistically being fair to women, we must look at the processors that increase precision and equal opportunity. Satisfying these fairness metrics would allow our model to predict more positives for women in an accurate way.

The reweighing technique decreases equal opportunity difference and increases precision slightly by 1%. It is the only metric which increases precision. The meta fair classifier improves the equal opportunity difference the most out of all the processors. In fact, the ratio becomes positive meaning that the model results in more accurately predicted women with access to emergency funds than accurately predicted men. Only looking at equal opportunity difference, this processor seems promising, but once we consider precision, we see a different picture. This processor decreased our precision from 68% to 37%. Since precision measures the proportion of true positives to all positives, this means we have increased our false positives. This makes sense because if our model increases the number of positives for women, there will be more true positives which will satisfy our equal opportunity difference fairness metric, but there will also be more false positives which will decrease precision.

This is a good example of how fairness can be complicated. We cannot just look at the processor that increases fairness across the board, but rather we must find a processor that prioritizes whatever the situation deems as fair. In our case fairness means something different than it would for an algorithm used in hiring practices. In a hiring scenario it would be a good idea to predict more women as being qualified for a job even if the false positive rate increases, because the model would be fighting existing prejudices and give women more opportunities. In our case, we want a model that classifies women just as accurately as it does men. We focused on minimizing false positive because we decided

it would be more harmful to classify someone as having access to emergency funds when in reality they do not. This interpretation might change depending on the context.

The equal opportunity difference is the most helpful metric in assessing accurately predicted positives. The processors that account for the fairness we care about are the reweighing, exponential gradient reduction, grid search reduction, and the reject option classifier. These processors maintain precision while increasing positive classifications for women.

Although these processors increase accurate positive outcomes for women, we lose some other kinds of fairness. The generalized entropy error, generalized entropy error difference, and between group generalized entropy error increase when the exponential gradient reduction, grid search reduction, and the reject option classifier processors are used. This means that individual fairness decreases overall. The generalized entropy error difference becomes negative meaning that men have more generalized entropy error than women. The between group generalized entropy error increased slightly meaning that there is more within group unfairness than between. These results are to be expected because group fairness and individual fairness metrics are contradictory. As we account for group fairness, we lose individual fairness. For access to emergency funds, we care more if our model is systematically being unfair to women, so it is okay that we lose some individual fairness.

Our study builds from numerous other studies that developed fairness metrics [26,27], developed processing techniques [14,21–23], and analyzed the trade-offs for different fairness metrics [26]. Our research expands on prior research through serving as an example of how these fairness metrics should be implemented. There must be considerations about the data, the model, and the metrics to decide how to process the model in the “fairest” way. Just knowing how the different fairness metrics are calculated is not enough. To ensure fairness one must understand how adjusting for different fairness metrics will impact the people that the model will be used on. If the wrong fairness metric is accounted for, the model might reinforce social injustices.

In our study, we only looked at a decision tree classifier. For future studies, implementing fairness metrics and processors in a specific context on multiple different machine learning algorithms could be helpful in seeing how these processors act for different types of models. There are other processors and fairness metrics that were not as applicable to our context that could be implemented in a future study. For example, a study implementing fairness metrics in a context where individual fairness needs to be prioritized. In our case, the model we built will be used primarily for educational purposes. Knowing the purpose of a model will aid in selecting the most relevant fairness metrics and processors.

Overall, the reweighing and exponential gradient reduction techniques are the most helpful processors for our decision tree classifier for emergency funds in Sub-Saharan Africa. The equal opportunity difference is the most informative fairness metric. For other models and other contexts, this may not be true. All of these processors account for fairness. However, it is important to find the processors that account for the specific type of fairness most relevant to the context and use of the model.

We plan to use this analysis as the basis of a workshop designed to teach undergraduate data science and computer science students how to effectively implement fairness and human rights into the machine learning workflow. We hope this work will serve as a useful comprehensive example for these students to see a human rights-based approach to machine learning in practice as well as a guide for students to integrate this approach into their own work. We seek to spread these ideas broadly with the long-term vision that all data scientists become aware of and embody the importance of centering ethics in machine learning and that this human rights-based approach becomes standard practice across the industry.

5. Appendix

List of predictors included in model

The 43 variables included in the model are listed below. See the [data documentation](#) for complete definitions of each variable.




Figure 8. Education Count by Gender

List of Predictors: 'wpid_random', 'gender', 'emp_in', 'fin2', 'fin7', 'fin14a', 'fin14b', 'fin15', 'fin16', 'fin17a', 'fin17b', 'fin19', 'fin20', 'fin21', 'fin22a', 'fin22b', 'fin24', 'sent_remittances', 'fin28', 'fin30', 'fin32', 'fin37', 'fin38', 'fin42', 'mobileowner', 'fin48', 'account_fin', 'account_mob', 'borrowed', 'receive_wages', 'receive_transfers', 'receive_pension', 'receive_agriculture', 'pay_utilities', 'pay_cash', 'pop_adult', 'wgt', 'age', 'educ', 'inc_q', 'remittances', 'access_cat', 'account_saved'

6. Acknowledgements

We would like to thank our sponsor organization [Women at the Table](#), and especially our supervisors/mentors Sofia Kypraiou and Caitlin Kraft-Buchman for creating the opportunity for us to do this project as well as for offering us guidance and feedback throughout the entire process. We would also like to thank our technical advisor, Dr. Shiya Cao, for providing extremely helpful Python resources and helping us work through challenges that we were having with our code. We would like to thank our Capstone professor, Dr. Albert Y. Kim, for providing conceptual and logistical guidance, especially in relation to machine learning knowledge, as well as for providing thorough feedback on our progress at various stages in the process. Lastly, we want to acknowledge Megan Lyster, Assistant Director of the Smith College Wurtele Center, for offering guidance on effective team collaboration and helping us to continually reflect on and adapt our process as a group throughout the course of the project.

Conflicts of Interest: The authors declare no conflict of interest.




Figure 9. Reasons for No Financial Account

References

1. Navarro, C.L.A.; Damen, J.A.; Takada, T.; Nijman, S.W.; Dhiman, P.; Ma, J.; Collins, G.S.; Bajpai, R.; Riley, R.D.; Moons, K.G.; others. Risk of bias in studies on prediction models developed using supervised machine learning techniques: systematic review. *bmj* **2021**, 375.
2. Hellström, T.; Dignum, V.; Bensch, S. Bias in Machine Learning—What is it Good for? *arXiv preprint arXiv:2004.00686* **2020**.
3. Barocas, S.; Hardt, M.; Narayanan, A. Fairness and machine learning.
4. Kleinberg, J.; Mullainathan, S.; Raghavan, M. Inherent trade-offs in the fair determination of risk scores. *arXiv preprint arXiv:1609.05807* **2016**.
5. Kypraiou, S. What is Fairness? **2021**. Publisher: PubPub.
6. Green, B.; Hu, L. The myth in the methodology: Towards a recontextualization of fairness in machine learning. Proceedings of the machine learning: the debates workshop, 2018.
7. Deho, O.B.; Zhan, C.; Li, J.; Liu, J.; Liu, L.; Duy Le, T. How do the existing fairness metrics and unfairness mitigation algorithms contribute to ethical learning analytics? *British Journal of Educational Technology* **2022**.
8. Kim, J.Y.; Cho, S.B. An Information Theoretic Approach to Reducing Algorithmic Bias for Machine Learning. *Neurocomputing* **2022**.
9. Anahideh, H.; Asudeh, A.; Thirumuruganathan, S. Fair active learning. *Expert Systems with Applications* **2022**, 199, 116981.
10. Nguyen, Q.H.; Ly, H.B.; Ho, L.S.; Al-Ansari, N.; Le, H.V.; Tran, V.Q.; Prakash, I.; Pham, B.T. Influence of data splitting on performance of machine learning models in prediction of shear strength of soil. *Mathematical Problems in Engineering* **2021**, 2021.
11. Juba, B.; Le, H.S. Precision-recall versus accuracy and the role of large data sets. Proceedings of the AAAI conference on artificial intelligence, 2019, Vol. 33, pp. 4039–4048.
12. Gupta, A.; Anand, A.; Hasija, Y. Recall-based Machine Learning approach for early detection of Cervical Cancer. 2021 6th International Conference for Convergence in Technology (I2CT). IEEE, 2021, pp. 1–5.
13. Zliobaite, I. On the relation between accuracy and fairness in binary classification. *arXiv preprint arXiv:1505.05723* **2015**.
14. Kamiran, F.; Calders, T. Data preprocessing techniques for classification without discrimination. *Knowledge and information systems* **2012**, 33, 1–33.
15. Menon, A.K.; Williamson, R.C. The cost of fairness in binary classification. Conference on Fairness, Accountability and Transparency. PMLR, 2018, pp. 107–118.
16. Zhou, N.; Zhang, Z.; Nair, V.N.; Singhal, H.; Chen, J. Bias, Fairness and Accountability with Artificial Intelligence and Machine Learning Algorithms. *International Statistical Review* **2022**.
17. Goodman, R. Why Amazon's Automated Hiring Tool discriminated against Women: News & Commentary, 2022.
18. Binns, R. On the apparent conflict between individual and group fairness. Proceedings of the 2020 conference on fairness, accountability, and transparency, 2020, pp. 514–524.
19. Mehrabi, N.; Morstatter, F.; Saxena, N.; Lerman, K.; Galstyan, A. A survey on bias and fairness in machine learning. *ACM Computing Surveys (CSUR)* **2021**, 54, 1–35.
20. Caton, S.; Haas, C. Fairness in machine learning: A survey. *arXiv preprint arXiv:2010.04053* **2020**.
21. Agarwal, A.; Beygelzimer, A.; Dudík, M.; Langford, J.; Wallach, H. A reductions approach to fair classification. International Conference on Machine Learning. PMLR, 2018, pp. 60–69.
22. Agarwal, A.; Dudík, M.; Wu, Z.S. Fair regression: Quantitative definitions and reduction-based algorithms. International Conference on Machine Learning. PMLR, 2019, pp. 120–129.
23. Pleiss, G.; Raghavan, M.; Wu, F.; Kleinberg, J.; Weinberger, K.Q. On fairness and calibration. *Advances in neural information processing systems* **2017**, 30.
24. Kamiran, F.; Karim, A.; Zhang, X. Decision theory for discrimination-aware classification. 2012 IEEE 12th International Conference on Data Mining. IEEE, 2012, pp. 924–929.
25. Celis, L.E.; Huang, L.; Keswani, V.; Vishnoi, N.K. Classification with fairness constraints: A meta-algorithm with provable guarantees. Proceedings of the conference on fairness, accountability, and transparency, 2019, pp. 319–328.

26. Speicher, T.; Heidari, H.; Grgic-Hlaca, N.; Gummadi, K.P.; Singla, A.; Weller, A.; Zafar, M.B. A unified approach to quantifying algorithmic unfairness: Measuring individual & group unfairness via inequality indices. *Proceedings of the 24th ACM SIGKDD international conference on knowledge discovery & data mining*, 2018, pp. 2239–2248.
27. Wachter, S.; Mittelstadt, B.; Russell, C. Why fairness cannot be automated: Bridging the gap between EU non-discrimination law and AI. *Computer Law & Security Review* **2021**, *41*, 105567.

© 2022 by the authors.